

Machine Politics in Chicago's Menu Program*

John C. Ruf

November 25, 2023

Abstract

In numerous urban and political systems, the reality of patronage clashes with theoretical models that prioritize marginal voters in determining outcomes. This paper examines the role of patronage using novel data from Chicago's menu program, where city council members distribute public goods in their wards using a fixed budget. I start with the case of Alderman Bernie Stone whom citizens rumored of under-providing precincts where he lost. I find strong evidence of a spending shift towards his opponents from his supporters after his removal from office. Next, I look for a broader trend using two differences-in-differences approaches. I find no significant shift comparing incumbent aldermen who barely won and lost reelection. Another model comparing indicted aldermen to "entrenched" aldermen shows a notable spending shift, albeit the results are sensitive. Precincts that supported an indicted alderman saw a 1.14% decrease in spending, while those that did not saw a 2.59% increase, amounting to a total gap of \$56,000 each year.

Keywords: Menu Program, Aldermen, Infrastructure Maintenance, Local Politics, Spatial Distribution, Machine Politics

JEL codes: H41 H72 L98

*I thank Jonathan Dingel, Anthony Fowler, Thomas Malthouse, and the participants of the 2022 UChicago RP Workshop for their helpful comments and suggestions. jcruf@uchicago.edu

1 Introduction

Do local politicians use machine-style politics to drive public goods towards their supporters? Is spending bias exacerbated or curbed by local political competition? Because most municipalities allocate spending through collective decision-making bodies such as councils, it is usually difficult to attribute expenditures to individual politicians. However, one such program exists in the United States: Chicago’s Aldermanic Menu Program. Through this program, the Chicago Department of Transportation (CDOT) gives each of Chicago’s 50 aldermen a \$1.5 million budget, an eponymous “menu” of common expenditures, and a map of 311 complaints in their ward. Using these resources, aldermen allocate funds to infrastructure maintenance, parks, and other public goods within their wards, even if they are not on the “menu.”

This paper asks whether the menu funding towards each alderman’s supporters or detractors changes after the alderman leaves office. If so, what determines the magnitude of this change? To answer this, this paper uses a differences-in-differences design to compare the time trends of each supporting or opposing precinct before and after the alderman leaves office. There are two cases where parallel trends are likely to hold: in close elections where the incumbent wins by a small margin and in cases where the incumbent is indicted and forced to leave office. In the former, both wards where the aldermen just barely won and just barely lost should have similar spending trends. Furthermore, both winning and losing aldermen in these cases should have similar expectations of winning, so they should behave similarly in the run-up to the election. In the latter, the timing of the in-office indictment is uncertain, so that should preclude any anticipation.

This paper makes two contributions, firstly it publicly introduces this setting’s data, and secondly it finds cases that supports Dixit and Londregan (1996)’s theory of machine politics. Dixit and Londregan (1996) argue that the key to machine politics is for politicians to be more efficient at allocating public goods to their supporters than their opponents. If this is the case, then the municipality arrives at a “machine politics” equilibrium where the incumbent targets all public goods to their top supporters. If there is no particular efficiency towards supporters, the municipality goes to a “downsian” equilibrium where both the incumbent and the challenger promise to allocate public goods to the most marginal voters. The paper finds some evidence for this bifurcation in the menu program setting. I find evidence that Bernie Stone, who was long rumored to do this, did indeed allocate more spending to his supporters than his opponents. I find a large swing in fund allocation to the lowest and highest 20% of precincts by campaign contributions after Stone’s 2011 loss to Debra Silverstein. The bottom 20% precincts saw \$20,000 more in annual public goods, while the top 20% precincts experienced a decline nearly double of that. I look for a broader trend using two differences-in-differences approaches, one focusing on comparing wards where the incumbent barely won reelection to those where they barely lost, and another comparing indicted aldermen to other “entrenched” aldermen. I find no significant shift comparing incumbent aldermen who barely won and lost reelection. But the indictment design shows a notable spending shift, similar in magnitude to the Stone case study.

Section 2 describes the program’s history, its rules, and its 2017 audit. Section 3 describes

the public economics literature on public goods allocation and lobbying and the spatial and urban economics literature on infrastructure allocation and how this paper relates to both literatures. Section 4 describes how the menu program data was collected, its summary statistics, and depicts an aggregate map of menu spending across Chicago for the past decade. Section 5 presents descriptive data for the case study of Alderman Bernie Stone, rumored to favor supporters with menu program funds. Section 6 explains this paper’s empirical strategy to determine if the Stone case study can be generalized to other aldermen. Section 7 shows the results of said empirical strategy, which finds that aldermen removed or retired due to criminal indictments, which has similar results to the Stone case study. However, this does not generalize to competitive elections. The results from the indictment-based differences-in-differences design are statistically weak and sensitive to the number of precincts included, but are economically large and significant. Section 8 concludes.

2 Background

The aldermanic menu program was initiated in 1994 and continues to this day (Ferguson, 2017). The program delegates approximately \$75 million every year to be split equally amongst the 50 aldermen in Chicago’s city council to be spent on projects they unilaterally select for their ward, given a “menu” of acceptable expenditures. Each alderman allocates approximately \$1.5 million per year to their ward. Chicago’s Office of Budget and Management (OBM) tracks the spending and its Department of Transportation (CDOT) manages the program. Each ward is defined to be approximately equal by population every 10 years, depending on the decennial census results, but the overall map is subject to a city council vote for approval. The Chicago Board of Elections draws the precinct map using the ward map, according to the available number of polling places, and each contain between 500 to 800 voters (Crowley, 2022). Each year in the spring, the Mayor, CDOT, and OBM send letters to the aldermen explaining the program and providing a menu of cost estimates and a list of possible projects. Before the aldermen select projects, CDOT and OBM provide a briefing packet with 311 complaint data. Finally, aldermen more or less spend the money as they see fit, with the only hard restriction being that they cannot spend more than \$1.32 million on any one project. Because this program and elections occur in the spring, new aldermen can only spend their entire menu funds the following year and often rely on the previous Alderman’s programmed budget the year they are elected.

“Off-menu” expenditures are also allowed, of which most “off-menu” funds go towards Parks, Chicago Public Schools, and miscellaneous beautification projects such as trees, murals, decorative garbage cans, designer bike racks, and more (Ferguson, 2017). While on-menu items are typically also provided by other funding sources within Chicago’s Capital Improvement Program, off-menu items such as murals and statues are usually directly credited to the aldermen, thus giving an incentive to reward supporters’ loyalty with public goods. The program is unique insofar as it gives elected politicians a wide berth over a significant portion of the City’s infrastructure budget and allows its use for items one does not typically think of as core infrastructure. An example of a

portion of a menu from 2013 is shown below in Figure 1.

Figure 1: An Example of a menu from 2012/2013

DEPARTMENT/PROGRAM	2012	2013
CDOT	ESTIMATED PRICING	ESTIMATED PRICING
Residential Street Resurfacing	\$36,000 for First 5 Blocks \$58,000 for Subsequent Blocks	\$38,000 for First 5 Blocks \$66,500 for Subsequent Blocks
Residential Alley Resurfacing	\$28,000 for First Alley \$37,500 for Subsequent Alleys	\$29,500 for First Alley \$47,500 for Subsequent Alleys
Green Alley Program	\$120,000 per Block	\$150,000 per Block
Alley Speed Hump Program	\$1,400 per Block	\$1,400 per Block
Concrete Alley Aprons	\$10,000 per Location	\$10,000 per Location
Street Speed Hump Program	\$3,700 per Block	\$3,700 per Block

The Chicago Office of the Inspector General audited the program in 2017 and found that the program was rife with misallocation – because wards are defined to be approximately equal by population but not equal by infrastructure needs, so some dense wards get 80% or more of their needs met while others get only 10% or less (Ferguson, 2017).

Thus, the OIG found that the program resulted in significant funding disparities between wards relative to infrastructure needs. Secondly, the OIG audit found that from 2012 through 2015, the program permitted aldermen to use \$ 15.1 million in menu funds for projects unrelated to so-called “core” infrastructure. Finally, the OIG audit found that CDOT allowed aldermen to use \$825,292 of menu funds on projects outside of the ward they were elected to represent so that they could spend it on the wards they were running for reelection in.

3 Literature Review

Most literature on political allocation of public goods focuses on theoretical models to explain behavior but comparatively little empirical testing of that behavior. This literature looks primarily at the political incentives to redistribute. The early literature in political economy focused on variations of targeting the median voter, the seminal paper being Downs (1957). This earlier literature was largely successful in explaining platform convergence in two-party systems. However, many political scientists and historians found this model wanting in addressing obvious cases of machine politics (Rakove, 1975; Golway, 2014). The harsh examples of the New York City Tammany Hall machine and the Chicago Democratic Machine are puzzling in a Downsian framework, as they maintained power through tight control of patronage jobs and distributing public goods to reward supporters. The resolution to this puzzle came in Dixit and Londregan (1996), which develops a general model of politics that encompasses both patronage and targeting median voters . The model

they use matches the Chicago setting to a tee: they assume that electorally-motivated politicians have a fixed budget to allocate and can distribute it a fixed number of pressure groups. In that model they find two equilibria: one where politicians allocate to supporters and one to marginal voters, and the equilibrium depends if the incumbent party is more efficient at distributing to supporters.

Since early contributions to the political economy literature Grossman and Helpman (1996), much work still focuses on identifying the effect of special interest groups on policy outcomes (Bombardini and Trebbi, 2020). This literature has taken different methodological approaches to quantifying this influence. One approach is to use structural models simulate outcomes. This previous literature finds that influencing politics has high returns, but that political influence has relatively small effects on policy outcomes on a national level and moderate effects at a local one (Kang, 2015; Finan and Mazzocco, 2021) Another approach is to take a more reduced-form approach and find some form of shock to identify how various outcomes change. One example, Margaret Frank, Hoopes and Lester (2022) find that governors in the US select place-based tax incentive locations more often when the tract’s state representative is a member of the governor’s party and is greatest with Republican governors. Another, Fowler, Garro and Spenkuch (2020) , uses a combination of regression discontinuity and first-differences design using shifts in betting market beliefs and finds that having a firm’s supported candidate win has little influence on the firms stock price.

This study is different from traditional political economy studies as it focuses on a municipal environment where the public good at play is primarily infrastructure. Thus, it is also related to the literature on municipal infrastructure provision and the new quantitative spatial economics literature. This includes Glaeser and Ponzetto (2018), Fajgelbaum et al. (2023), and Bordeu (2023). Glaeser’s seminal paper focuses on infrastructure’s “visible” and “invisible” costs. Overall voters prefer overspending, as see the benefits of infrastructure but not the costs. Local voters prefer under-building, as they see the externalities of the infrastructure. Glaeser uses this framework to explain why urban mega-projects declined in the US, as urban voters became more organized. Fajgelbaum et al. (2023) examine how political economy influenced the planning of California’s high-speed rail (CHSR) project. They find that preferences for widespread approval lead to the planner placing CHSR stations farther from dense metro areas than a politically blind planner. Finally, Bordeu (2023) looks at how infrastructure is allocated across a similarly decentralized city, Santiago, Chile, and finds that the sub-city municipalities over-invest in core areas and under-invest in areas near their boundary using a quantitative spatial model. She finds that infrastructure centralization would increase aggregate infrastructure investment and yield large welfare gains.

4 Data Description

The primary data this paper employs are located menu expenditures, which contain the yearly allotted allocations from aldermen and their respective locations from 2005 to 2022. This dataset

comes from menu spending reports that are publicly available from 2011 through 2022, and records that were not previously publicly available that were obtained through a FOIA request to the OBM OBM (2022). I then scraped these PDFs and used the resulting cost total, ward, and location description data to create a map of shapes of the locations of the expenditures. I then used the location description text to locate each project’s described vertices using the Census geocoding API. If the Census’s API failed, I automatically used Google Maps’ API instead. In total, 43,596 projects needed to be located, and 83% of them were successfully located using one of the above methods. For example, spending on playground equipment would be a singular point, while spending on a street would be a line, and spending on all alleys within a given block would be a polygon. This dataset contains 41,381 precinct-year observations, which record the total spending in a given voting precinct in a given year.

Figure 2 depicts two side-by-side histograms of the distribution of spending per precinct aggregated across the 2005-2011 period, which used the 2003-2011 ward boundaries and the 2012-2022 period with the 2012-2022 ward boundaries. The decentralized nature of the menu program leads to a considerable variation in spending per precinct, but the distribution has a long right tail. Both figures are winsorized at the 99th percentile to remove outliers.




Figure 2: Distribution of Spending per Precinct for both ward maps in the dataset

Next, Figure 3 depicts within-ward spending variation across Chicago using the 2012-2022 precinct map; spending is often concentrated in a few precincts.




Figure 3: Map of Spending per Precinct, 2012-2022

To avoid issues with different levels of “observable” spending from year to year, I use the fraction of spending located in a precinct in a given year as the dependent variable for subsequent analyses. “Observed” means one of the two methods above successfully located it to a location in Chicago. Thus, if Y_{py} is the fraction of spending located in precinct p in year y , then $Y_{py} = \frac{S_{py}}{S_{wy}}$, where S_{py} is the spending in precinct p in year y , and S_{wy} is the total amount of ward w ’s successfully located spending in year y . Table 1 depicts summary statistics for this variable. Due to the long right tail of the distribution, the mean is much larger than the median. Half of precincts get almost no spending.

Table 1: Summary statistics of fraction of total spending by precinct from 2012 to 2022

mean	median	sd	upper quartile	lower quartile
2.42	0.23	4.48	3.35	0

TODO: Add a scatterplot of spending fraction per precinct vs. net votes for the incumbent alderman

5 Bernie Stone Case Study

This section discusses the Bernie Stone case study, which is the primary motivation for this paper. Bernie Stone was an alderman in Chicago’s 50th ward from 1973 to 2011. He was well known for


his, “political philosophy.”

“You take care of the people who take care of you — you know, the people who voted for you, That’s not Chicago politics, that’s Politics 101.” - Alderman Bernie Stone (50th ward), Quoted from Zekman (2009)

In fact an alderman who grew up in the 50th ward once remarked that,

“Well, I grew up in the 50th Ward and you know, God bless [the late former Ald.] Bernie Stone, may he rest in peace, but I remember crossing California going west, every street was resurfaced almost every year. They always had brand new lighting and then east of California, where he would lose the precincts consistently, I mean the streets were in shambles. Many people felt he was spending the bulk of the menu money west of California, where he was getting the bulk of the vote.” - Alderman Carlos Ramirez-Rosa (35th Ward), Quoted from Vevea (2017)

This quote is a clear example of the type of behavior that this paper seeks to investigate. This phenomenon could not be quantified previously because the CDOT did not make the 2005-2011 documents publicly available. Furthermore, the 2011-2022 data was in PDF form, making locating the spending tedious and difficult. This paper is the first to put numbers to this anecdotal evidence. Figure 4 depicts the precincts that supported Stone in the 2007 runoff election and the precincts that gave Stone the most individual contributions. Combined, both maps show that the southwestern portion of the ward is the most supportive of Stone, on average.



(a) Campaign contributions to Alderman Stone, 2003-2011

(b) Net votes for Alderman Stone, 2007

Figure 4: Alderman Stone’s support in the 50th ward

Figure 5 gathers the top and bottom quintile of 44 precincts in the 50th ward by contributions to Stone and shows the average fraction of the total located menu budget spent in each quintile and the rest of the ward. Note that the precincts that contributed the most to Stone’s campaign are the same precincts that contributed the most net votes to Stone in the 2007 runoff election. After his reelection in 2007, Stone’s spending per precinct was heavily concentrated in the precincts that supported him. After his defeat, spending per precinct is much more evenly distributed across the ward.


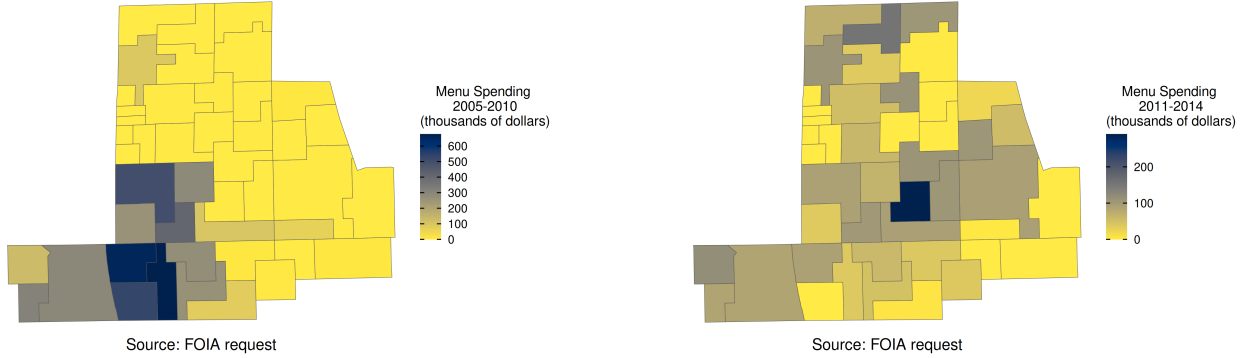


Figure 5: Average spending by precinct group in the 50th ward, 2005-2016

Figure 6 shows the shift geographically. It depicts a clear movement away from the south-western portion of the ward to a roughly even distribution after Stone’s defeat.



(a) 50th ward menu allocation, 2005-2010

(b) 50th ward menu allocation, 2011-2015

Figure 6: 50th ward menu maps before and after Stone’s defeat

6 Empirical Framework

This section discusses the empirical framework used to analyze the data. I use a differences-in-differences approach to determine whether aldermen are allocating spending to precincts that support them. However, I focus on a heterogeneous treatment effect robust estimator. The most standard specification is as follows:

$$Y_{pt} = \alpha_p + \gamma_t + \beta T_{pt} + \epsilon_{pt} \quad (1)$$

Y_{pt} is the fraction of observed spending in precinct p in year t . T_{pt} is a dummy variable equal to 1 if the incumbent alderman was removed from office in year t . α_p and γ_t are precinct and year fixed effects respectively. ϵ_{pt} is the error term. β represents how much spending on an average precinct in the sample changes after an alderman is removed from office. Under parallel trends, this represents the causal impact of removing an alderman from office on spending in the precincts contained in the sample. Thus, in the case of primarily supporting precincts, a negative β would indicate that the incumbent alderman was allocating more spending to those that supported them than the following alderman would have chosen. Yet, this model should not be naively applied with two-way fixed effects due to the recent literature on heterogeneous treatment effects in differences-in-differences designs. The growing literature shows that the standard two-way fixed effects estimator biases the coefficient β , is a weighted average of several differences-in-differences that compare how Y_{pt} progress across pairs of groups, of which some of them can be treated in both periods (de Chaisemartin and D’Haultfœuille, 2020) (Callaway, Goodman-Bacon and Sant’Anna, 2021). I use the heterogeneous-treatment effect robust estimator proposed by (Callaway and Sant’Anna,

2021) to address this issue.

Why not use a continuous treatment design, with an indicator of support as a marker of treatment strength? The reason for this comes from Callaway, Goodman-Bacon and Sant’Anna (2021). A much stronger form of parallel trends must be satisfied for a continuous treatment design to be valid. Callaway, Goodman-Bacon and Sant’Anna (2021) describe the set of assumptions that allow the identification of the “average causal response of treatment (ACRT) function” that is a weighted average of causal responses to a unit change in treatment intensity. They find that identifying an ACRT function requires not only the traditional, support, common trends, and no anticipation assumptions, but also a “homogenous ATT” assumption. This assumption is necessary because of the selection bias induced when treatment groups who elect into stronger treatment do so because the stronger treatment to benefits them more than other groups. This assumption means politicians only allocate spending based on support, regardless of precinct characteristics. That is, no precincts are more efficient at giving support than others. However, this assumption is directly at odds with the Dixit and Londregan model, which shows that for machine politics to exist, the politician must be able to distribute goods more efficiently to supporters than non-supporters. This assumption is just the converse of precincts being more efficient at organizing support to the incumbent than others. Thus, using a continuous treatment design would be effectively equivalent to assuming that machine politics does not exist in this context.

This paper estimates four variations of Equation 1. The first two rely on a close-election assumption to justify the parallel trends assumption. This assumption means that incumbent aldermen who win by a small margin have similar spending trends to those who lose by a small margin. This study defines a close margin as 10% or less; this corresponds to approximately within 1,200 votes. Therefore, they should behave similarly in the run-up to the election as they have similar expectations of winning. We then examine two sets of precincts. The first set of precincts is the top quintile (8) precincts by vote margin for each ward in the 2015 or 2019 election, whichever is appropriate to the treatment group. The second set of precincts is the bottom quintile precincts by vote margin for each ward in the 2015 or 2019 election, whichever is appropriate to the treatment group. We use the 2012 through 2022 years to estimate this model, due to the 2011 redistricting complicating the use of 2005-2011 data.

The second two variations rely on a simultaneous set of indictments of aldermen in 2019, causing three aldermen to either be ineligible for reelection or to retire. These aldermen were Daniel Solis, Ricardo Munoz, and Willie Cochran. Daniel Solis left office after being caught by the FBI and wearing a wire to record Alderman Ed Burke, who was indicted in 2019 but not removed from office until 2023. Ricardo Munoz retired shortly after reporters discovered that he had spent PAC money on personal expenses. Willie Cochran retired after pleading guilty to wire fraud and misusing campaign funds for gambling and personal expenses. We compare this group to a set of 10 control aldermen who were not indicted, have been in office for at least 10 years, and won reelection in 2019 in the general election, indicating they were not in a competitive election. To measure whether or not a precinct supported an alderman, I use the total number of campaign contributions donated

to the aldermen from the precinct in the 2015 and 2019 elections. We do this because many of the “entrenched” aldermen have not faced a competitive election in decades, so net votes are not a good measure of support. Despite this, all aldermen still accept campaign contributions even when there is no challenger, so I use this to measure support. In this case, I expect the trends for the indicted and control aldermen will be the same, as they are all not in competitive elections and thus their spending preferences should be stable over time. Furthermore, the unexpected timing of indictments within the election would hopefully preclude any anticipatory behavior. In both cases, I cluster standard errors at the ward level, as that is the level at which the treatment is assigned.

7 Results

I compute the dynamic aggregation of the average treatment effect on the treated group for the least and most supporting precincts and display the results in Table 2. Also shown are the estimate’s standard errors, confidence intervals, and Wald-test p-values for the pre-test of the parallel trends assumption. I see that both estimated ATTs are not statistically significant for the close election design. Furthermore, the least supporting precinct’s ATT is not economically significant. Half a percent of the menu program’s budget is roughly \$7,500. That is not even enough to afford two street-wide speed bumps. Additionally concerning are the extremely low p-values for the pre-trends test. This test indicates that the close election design likely does not guarantee that parallel trends hold. This failure may be because factors influencing the trend of infrastructure spending and needs correlate with the election outcome. Alternatively, it could be because the margin chosen is too large to ensure that the two groups are similar. The results of the competitive elections design is susceptible to the number of precincts included, often changing the sign of both estimated ATTs. Next, I can look at the results of the indictment design in. I note that the treatment effect for the least supporting precincts is over five times larger than the treatment effect estimated in the competitive election design. I also see a much larger Wald-test p-value for the pre-trends test. These values indicate that perhaps the indictment design’s assumption is more believable than the close election design’s assumption. However, I also see that the most supporting precinct’s treatment effect, while statistically significant, is similar in magnitude to the treatment effect of the competitive election design. The results shown Table 2 are only sensitive to increasing the number of precincts, not decreasing it.

Table 2: Comparison of average treatment effects

	Competitive Election		Indictment	
	Opposing ATT	Supporting ATT	Opposing ATT	Supporting ATT
ATT	0.50	-1.51	2.59	-1.15
Std. Error	0.45	1.12	0.78	0.32
95% Conf. Int.	(-0.38, 1.39)	(-3.71, 0.68)	(1.06, 4.11)	(-1.77, -0.52)
Pre-Trends P-value	0.005	0.076	0.199	0.174
Obs	1680	1680	1144	1144

TODO: Add a figure that shows how the ATT of each changes for each design as the number of precincts included changes from 5 to 12.

I can look at how the average treatment effect changes over time in Figure 7, which depicts how the least supporting precincts estimated treatment varies over time in the close election context. As inferred from the standard errors in Table 2, the estimated effect over time is noisy. This is likely because if you think about an optimal social planner, they will not allocate spending using a highly auto-correlated spending rule. More likely, they would allocate spending based on the state of decay of the infrastructure stock of the precincts. If you imagine that close elections push aldermen closer to a hypothetical social planner, you would expect lots of noise in the estimated effect.




Figure 7: Average treatment effect over time for least supporting precincts: competitive election design

Finally, I look at how the average treatment effect changes over time in Figure 8, which depicts how the least supporting precincts estimated treatment varies over time in the indictment design. The standard errors, while large, are much smaller, and the mean estimate is much more stable. Figure 8 also shows evidence of anticipation in the year before the first treatment, as the 2019 estimate is significantly smaller than the other estimates. This holds for the most supporting precincts as well, albeit in reverse. Thus, as the incumbent alderman leaves office, they allocate more spending to their most supporting precincts than their least keeping precincts. These results must be taken with a grain of salt due to the high standard errors and the likely pre-trends violation.




Figure 8: Average treatment effect over time for least supporting precincts: indictment design

TODO: Create custom figures because the did's package figures are ugly.

8 Conclusions

Overall this paper finds some evidence that aldermen may sometimes disproportionately allocate spending to their most supporting precincts, particularly when they are long-entrenched and not facing a competitive election. The paper starts by verifying a rumor that a particular alderman disproportionately allocated spending to his most supporting precincts. Then it uses two applications of a differences-in-differences research design to arrive at this fact. The first application focuses on aldermen who lost by a small margin, and finds that the evidence for disproportionate spending is very weak. While the magnitude of the estimated effect on the top precincts in the close election design is large, it is not even close to statistically significant, and even the sign of the effect is not robust to changes in the number of supporting/opposing precincts included in the sample. I find that the effect is economically large and statistically significant when under the indictment design. The statistical significance is sensitive to parameters such as the number of supporting or opposing precincts included per ward, but the economic significance is not and stays largely the same for this design so long as you decrease the number of precincts. This is likely due to the fact that increasing the number of precincts dilutes the average treatment effect, so it decreases the aggregated treatment effect.

The results help build on the burgeoning urban economics of infrastructure literature by show-

ing that political incentives can distort the allocation of infrastructure spending. Secondly, the differences between the competitive election and indictment designs show that the electoral competition can be a powerful force in constraining the clientelistic tendencies of politicians. There is also a lesson in urban planners that while discretion can be useful, it can also be abused and lead to unintended consequences. Therefore, the capacity for discretion should be carefully considered when designing a program.

References

- Bombardini, Matilde, and Francesco Trebbi. 2020. “Empirical Models of Lobbying.” Annual Review of Economics, 12(1): 391–413.
- Bordeu, Olivia. 2023. “Commuting Infrastructure in Fragmented Cities.” University of Chicago.
- Callaway, Brantly, and Pedro H.C. Sant’Anna. 2021. “Difference-in-Differences with multiple time periods.” Journal of Econometrics, 225(2): 200–230. Themed Issue: Treatment Effect 1.
- Callaway, Brantly, Andrew Goodman-Bacon, and Pedro H. C. Sant’Anna. 2021. “Difference-in-Differences with a Continuous Treatment.”
- Crowley, Kinsey. 2022. “Chicago’s cutting precincts just before November election.”
- de Chaisemartin, Clément, and Xavier D’Haultfoeuille. 2020. “Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects.” American Economic Review, 110(9): 2964–96.
- Dixit, Avinash, and John Londregan. 1996. “The Determinants of Success of Special Interests in Redistributive Politics.” The Journal of Politics, 58(4): 1132–1155.
- Downs, A. 1957. An Economic Theory of Democracy. Harper.
- Fajgelbaum, Pablo D, Cecile Gaubert, Nicole Gorton, Eduardo Morales, and Edouard Schaal. 2023. “Political Preferences and the Spatial Distribution of Infrastructure: Evidence from California’s High-Speed Rail.” National Bureau of Economic Research Working Paper 31438.
- Ferguson, Joseph. 2017. “CHICAGO DEPARTMENT OF TRANSPORTATION ALDERMANIC MENU PROGRAM AUDIT.” Office of Inspector General.
- Finan, Frederico, and Maurizio Mazzocco. 2021. “Electoral incentives and the allocation of public funds.” Journal of the European Economic Association, 19(5): 2467–2512.
- Fowler, Anthony, Haritz Garro, and Jörg L. Spenkuch. 2020. “Quid Pro Quo? Corporate Returns to Campaign Contributions.” The Journal of Politics, 82(3): 844–858.
- Glaeser, Edward L., and Giacomo A.M. Ponzetto. 2018. “The political economy of transportation investment.” Economics of Transportation, 13: 4–26. The political economy of transport decisions.

- Golway, T. 2014. Machine Made: Tammany Hall And The Creation Of Modern American Politics. WW Norton.
- Grossman, Gene M., and Elhanan Helpman. 1996. "Electoral Competition and Special Interest Politics." The Review of Economic Studies, 63(2): 265–286.
- Kang, Karam. 2015. "Policy Influence and Private Returns from Lobbying in the Energy Sector." The Review of Economic Studies, 83(1): 269–305.
- Margaret Frank, Mary, Jeffrey L. Hoopes, and Rebecca Lester. 2022. "What determines where opportunity knocks? Political affiliation in the selection of Opportunity Zones." Journal of Public Economics, 206: 104588.
- OBM, Chicago. 2022. "Capital Improvement Archive." Website.
- Rakove, M.L. 1975. Don't Make No Waves—don't Back No Losers: An Insider's Analysis of the Daley Machine. A Midland Giant, Indiana University Press.
- Vevea, Becky. 2017. "What Would You Do With A Million Dollars? Whether Participatory Budgeting Is Worth The Effort." WBEZ Chicago.
- Zekman, Pam. 2009. "Better Government Association: How Do Aldermen Spend Their 'Piggy Bank' Money?" Website.